

## Week 2 :

### Resources

- **Code and Tutorials:**
  - **Pre-trained CNN Models:**
    - **PyTorch Models:** <https://pytorch.org/vision/stable/models.html>
    - **TensorFlow Models:** [https://www.tensorflow.org/api\\_docs/python/tf/keras/applications](https://www.tensorflow.org/api_docs/python/tf/keras/applications)
  - **Transformer Tutorials:**
    - **Hugging Face Transformers:** <https://huggingface.co/transformers/>

### Tasks:

#### 1. Data Preparation

- **Identify and Download a Dataset:** download dataset given in link below
- **Preprocessing Captions:**
  - Tokenize captions using NLP libraries (e.g., NLTK, SpaCy, Hugging Face).
  - Build a vocabulary and encode captions into sequences.
  - Add padding/truncation for captions to ensure uniform length.
- **Preprocessing Images:**
  - Use a pre-trained CNN (e.g., ResNet, EfficientNet) to extract visual features from images.
  - Normalize and resize images to a fixed size (e.g., 224x224 for most CNNs).
  - Store extracted features for efficient training.

#### 2. Model Prototyping

- **CNN Feature Extractor:**
  - Implement a CNN (like ResNet-50) to process images and extract visual embeddings.
- **Text Encoder (Transformer or LSTM):**
  - Use an encoder to process tokenized captions, experimenting with LSTMs or transformer-based models (e.g., Hugging Face's BERT).
- **Combine Image and Text Features:**
  - Add a mechanism to concatenate or fuse image embeddings and encoded captions for the decoder to process.
- **Initial Decoder Setup:**
  - Use a transformer decoder or an attention mechanism to generate the captions sequentially.
  - Begin by building a minimal model pipeline and confirm it can process dummy data (input image + captions) without errors.

**Assignment dataset :** <https://www.kaggle.com/datasets/adityajn105/flickr8k>

### Expected Outcomes:

- Preprocessed dataset with tokenized captions and extracted image features.

- A working model prototype combining CNNs and transformers for caption generation.

## HINTS

### Preprocessing Captions:

**Tokenization:** Use a tokenizer to split sentences into words:

```
from nltk.tokenize import word_tokenize
```

```
captions = ["A man riding a horse.", "A dog playing in the yard."]
```

```
tokenized_captions = [word_tokenize(caption.lower()) for caption in captions]
```

- **Build Vocabulary:** Assign unique indices to words. Include special tokens like `<start>`, `<end>`, and `<pad>`.

**Encode Captions:** Convert tokens into sequences of integers:

```
word_to_index = {"<start>": 1, "a": 2, "man": 3, ...} # Example vocabulary
```

```
encoded_caption = [word_to_index[word] for word in tokenized_caption]
```

- 

**Pad Captions:** Use libraries like TensorFlow/Keras for padding:

```
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
padded_captions = pad_sequences(encoded_captions, maxlen=20, padding="post")
```

- 

---

### Preprocessing Images:

Use pre-trained CNN models to extract features:

```

from tensorflow.keras.applications import ResNet50

from tensorflow.keras.applications.resnet import preprocess_input

from tensorflow.keras.preprocessing.image import img_to_array,
load_img

# Load and preprocess image

image = load_img('image.jpg', target_size=(224, 224))

image = preprocess_input(img_to_array(image))

# Extract features

model = ResNet50(weights='imagenet', include_top=False, pooling='avg')

features = model.predict(image.reshape(1, 224, 224, 3))

•

Save the extracted features for faster processing during training:
import pickle

with open('features.pkl', 'wb') as f:

    pickle.dump(features, f)

•

```

---

## 2. Model Prototyping

### CNN Feature Extractor:

- Use a CNN to encode image features. You can experiment with pre-trained networks like ResNet, EfficientNet, or MobileNet.
- The output of the CNN should be a fixed-length feature vector for each image.

---

### Text Encoder:

Use a transformer-based encoder or LSTM for captions:

```
from tensorflow.keras.layers import Embedding, LSTM

vocab_size = 5000 # Adjust based on your dataset

embedding_dim = 256

max_length = 20

text_input = Input(shape=(max_length,))

embedded_text = Embedding(vocab_size, embedding_dim)(text_input)

encoded_text = LSTM(512, return_sequences=True)(embedded_text)
```

- 

---

### Combine Image and Text Features:

Concatenate image features with encoded text:

```
from tensorflow.keras.layers import Concatenate

combined_features = Concatenate()([features, encoded_text])
```

- 

### Initial Decoder Setup:

- Use a transformer decoder or attention mechanism:
  - Implement cross-attention to align image features with text.

Keras has built-in layers for attention in TensorFlow 2.6+:

```
from tensorflow.keras.layers import Attention
```

```
attention = Attention()([combined_features, encoded_text])
```

-